

# Doctoral Project Presentation

## DEDS Winter School 2022

April 5, 2022

Abduvoris Abduvakhobov

Center for Data Intensive Systems, Daisy  
Department of Computer Science  
Aalborg University  
Denmark



**AALBORG UNIVERSITY**  
DENMARK

# Agenda



Introduction

Motivation

- Data and Paramount Properties

State of the Art

- Time Series Management Systems

- ModelarDB

- Open Issues

Project Objectives

Work and Publication Plan

- Work and Time Plan

- Tentative Publication List

Secondment

# Introduction



**PhD Topic:** ESR:2.3. Model-based storage for time series

**Supervisors:**

- Associate Professor Christian Thomsen, AAU
- Professor Esteban Zimanyi, ULB

**Secondment:** Siemens Gamesa Renewable Energy

**PhD Start Date:** February 1, 2022

**Current progress:**


- Finalizing doctoral project plan
- Started collaboration with secondment partner and got access to data

# Motivation

Research undertaken by Jensen et al. 2021

## Meetings with manufacturers, owners, and energy traders:

- Modern turbines are monitored by up to 7,000 high-quality sensors
- The storage needed makes storing high-frequency sensor data infeasible
- Simple aggregates (e.g. 10-minute averages) are stored instead of the high-frequent series, thereby removing useful fluctuations and outliers:



- Users believe problems can be found earlier with high-frequency data.
- Compression need only be lossless for some types of time series.

# Motivation

## Data



### Meetings with manufacturers, owners, and energy traders:

- The sensors are installed with wired power and connectivity.
- Each sensor produces a data stream sampled to, e.g., a 10 Hz series.
- Collected measures include: Air Pressure, Humidity, Voltage, Power, Rotation Speed, Temperature, Wind Direction, Wind Speed, Internal Controller Measurements.
- The time series are regular, cleaned, but gaps without values can occur.
- Metadata for each time series must also be stored, e.g., as dimensions.



# Motivation

## Paramount Properties

Paramount properties for a system managing wind turbine data:

**Distribution:** The system must be able to scale to many nodes.

**Stream Processing:** Data points are arriving continuously as a regular time series and must be queryable with a short latency.

**Compression:** High compression is needed for high-frequency data.

**Efficient Retrieval:** Indexes or ordered storage for fast retrieval.

**Approximate Query Processing:** Approximate answers can be accepted for some time series and enables use of lossy compression.

**Extensibility:** Allows users with domain knowledge to implement new storage methods optimized specifically for their data sets.

# State of the Art

## Time Series Management Systems



### Time Series Management System or TSDB:

- Stores time series that consists of time stamp.
- Optionally contain metadata or tags.
- Processes queries on time series.
- Queries contain timestamp or a time range.

### Categorization based on architecture:

- Internal Data Stores:
  - Mostly centralized.
  - Tightly coupled data storage to processing component.
  - Few mature implementations.
  - Examples: Plato, LittleTable, VergeDB, Chronos, Apache IoTDB

# State of the Art

## Time Series Management Systems



### Categorization based on architecture:

- External Data Stores:
  - Predominantly distributed.
  - New processing engine on top of external data store.
  - Most number of mature implementations.
  - Examples: Apache Druid, Bolt, Gorilla, BTrDB and ModelarDB
- Extension for RDMS:
  - Extends popular RDBMS.
  - predominantly centralized.
  - Small number of mature implementations.
  - Examples: Chronix, EdgeDB, and Heracles






# State of the Art


## ModelarDB

### ModelarDB

- Individual time series can be described with models:
- E.g.,  $v = a * t + b$  can represent a sub-sequence using only  $a$  and  $b$ .




- Uses Apache Cassandra for storage and Apache Spark for query processing.
- Approximates the time series values using mathematical functions (models) and stores only model coefficients.
- Currently includes three different model types:



# ModelarDB

## Correlated Time Series

- A data set often contains redundant information across time series:
  - E.g, co-located temperature sensors often produce similar values.
- ModelarDB can group correlated time series together and compress them as one stream of models to reduce the storage required.
- A list of model types fit models to data points, e.g., a constant (PMC-Mean), linear (Swing), and lossless (Gorilla) model type:



(Jensen et al. 2021)

# ModelarDB


## Multi-model Group Compression



- Time series are grouped based on user hints given using primitives.
- The primitives can be combined and allow users to state that series are correlated based on their source or their dimensional hierarchy.
- Users can use their domain knowledge, analyze historical data, or use ModelarDB's automatic grouping method built on the primitives.

### Grouping 9632 and 9634:

- From specific sources: 9632 and 9634
- Sharing a specific member:  
Location 3 Aalborg
- Share members until a level:  
Location 3
- The dimension's distance: 0.25
- Automatically (distance): auto



# ModelarDB

## Open Issues



## Open Issues

- No integrated functionality for evaluating the efficiency and use of model types;
- Model types are quite generic;
- Sampling interval and error bound could be changed dynamically;
- System only supports automatic and manual grouping using heuristics (domain knowledge and metadata);
- Time series must be ordered and have a regular sampling interval;

# Project Objectives

## Research Questions



- RQ1: How can we efficiently evaluate the compression performance of model types and the quality of compression to varying error bounds of ModelarDB on different datasets?
- RQ2: Depending on the the outcomes of RQ1, what other model types can be implemented to improve the compression and query performance of ModelarDB on real-life RES datasets?
- RQ3: How can time series automatically be grouped using different correlation statistics and provided heuristics during the ingestion process?
- RQ4: How can model-based ingestion of time series with a dynamic sampling interval and error bound be supported in ModelarDB?



# Work and Publication Plan

## Work and Time Plan

Time	Plan
Semester 1, Spring' 22 (Home)	Literature study Problem formulation Preparation of the Doctoral Project Plan Begin work on Paper 1 Begin development of ModelarDB performance evaluation tool Data collection Start general and project-related courses Trying out ModelarDB on real-life datasets and analyzing the results
Milestones	<b>Submission of 2-month Doctoral Project Plan</b> <b>Establish collaboration with secondment partner and get access to data</b>
Semester 2, Fall' 22 (Home)	Continue project-related courses Develop and test performance evaluation tool for ModelarDB
Milestones	<b>Submission of Paper 1</b> <b>Submission of 11-month Doctoral Project Plan</b> <b>Secondment completed</b>

# Work and Publication Plan

## Work and Time Plan



Time	Plan
Semester 3, Spring' 23 (Host)	Develop new model types for ModelarDB Begin work on Paper 2
<b>Milestones</b>	<b>Solution design of new model types</b>
Semester 4, Fall' 23 (Host)	Refine and test new model types with real-life datasets Begin work on Paper 3
<b>Milestones</b>	<b>Submission of Paper 2</b>
Semester 5, Spring' 24 (Home)	Develop new method for correlation-based grouping in ModelarDB Begin work on Paper 4
<b>Milestones</b>	<b>Submission of Paper 3</b> <b>Completed all general and project-related courses</b>
Semester 6, Fall' 24 (Home)	Develop and test new method for compressing dynamic sampling intervals and model error bounds Writing the Thesis
<b>Milestones</b>	<b>Submission of Paper 4</b> <b>Submission of PhD Thesis</b>

# Work and Publication Plan

## Tentative Publication List



**Tentative Title of Paper 1:** A tool for analysis of the efficiency of model-based compression in ModelarDB

**Type:** Conference paper.

**Description:**

- Analysis of the efficiency of current model types deployed by ModelarDB.
- Integrated tool that explains the system performance and its usage of model types.
- Performance indicators and visualization.

**Datasets:** Siemens Gamesa Renewable Energy and ENGIE data.

**Authors:** A. Abduvakhobov, S.K. Jensen, C. Thomsen, T. B. Pedersen, E. Zimányi, T. Pasma.

**Length:** 12 pages.

**Time of submission:** September, 2022.

**Outlet:** IEEE BigData.



# Work and Publication Plan

## Tentative Publication List



**Tentative Title of Paper 2:** New model types to achieve better compression rate and lower error bound for ModelarDB.

**Type:** Conference paper.

**Description:**

- Develops new model types for better ingestion and storage use.
- Mainly tailored to match real-life use cases.
- Novel time series compression method.

**Datasets:** Siemens Gamesa Renewable Energy and ENGIE data.

**Authors:** A. Abduvakhobov, S.K. Jensen, C. Thomsen, T. B. Pedersen, E. Zimányi, T. Pasma.

**Length:** 12 pages.

**Time of submission:** June, 2023.

**Outlet:** EDBT.

# Work and Publication Plan

## Tentative Publication List



**Tentative Title of Paper 3:** Automatic grouping of time series by deploying correlation statistics in ModelarDB

**Type:** Journal paper.

**Description:**

- More optimized and faster grouping (possibly in a streaming fashion).
- Leverages correlation and other statistical attributes of time series.
- Supported by user heuristics and metadata.

**Datasets:** Siemens Gamesa Renewable Energy and ENGIE data.

**Authors:** A. Abduvakhobov, S.K. Jensen, C. Thomsen, T. B. Pedersen, E. Zimányi, T. Pasma.

**Length:** 12 pages.

**Time of submission:** March, 2024.

**Outlet:** PVLDB.

# Work and Publication Plan

## Tentative Publication List



**Tentative Title of Paper 4:** Adding dynamic sampling intervals and error bounds for time series ingestion of ModelarDB.

**Type:** Conference paper.

**Description:**

- Dynamic error bound and sampling interval.
- User controls the error bound and sampling interval.
- More fine-grained data for exceptional cases.
- Also discusses automatic adjustment of error bound.

**Datasets:** Siemens Gamesa Renewable Energy and ENGIE data.

**Authors:** A. Abduvakhobov, S.K. Jensen, C. Thomsen, T. B. Pedersen, E. Zimányi, T. Pasma.

**Length:** 12 pages.

**Time of submission:** October, 2024.

**Outlet:** ICDE.

# Plan for PhD Courses



Course Name	At	Type	ECTS	Time	Status
General Courses	AAU	General	13.75	'22-'25	Planned
Project-related Courses	AAU	General	17	'22-'25	Planned
Winter School (ARC)	ARC	General	3	Spring'22	Mandatory
Summer School (ULB)	ULB	Project	3	Summer'22	Mandatory
Winter School (AAU)	AAU	General	3	Winter'22	Mandatory
Summer School (UPC)	UPC	Project	3	Summer'23	Mandatory
Conference Attendance	TBD	Project	3	TBD	Planned
Danish Lessons	TBD	General	TBD	TBD	Mandatory
<b>Total: 30.75</b>					

# Secondment

**Partner Organization:** Siemens Gamesa Renewable Energy

**Secondment Supervisor:** Tjip Pasma

**Secondment start date:** March 21, 2022

**Data:** power, voltage, reactive power, controller measurements, and other metadata

**Size:**  $\approx 110$  TB

Out[13]:

	TimeStamp	Component	Measured	...	...	...	...	...	...	...	...
0	2021-10-01 16:00:00.123	1	200.0	173.825897	165.654297	165.905594	165.905380	173.295059	-9.18	200.0	...
1	2021-10-01 16:00:00.273	1	200.0	173.824905	165.663910	165.903992	165.905212	173.291351	-9.18	200.0	...
2	2021-10-01 16:00:00.423	1	200.0	173.825897	165.673157	165.902466	165.905060	173.287689	-9.18	200.0	...
3	2021-10-01 16:00:00.570	1	200.0	173.821899	165.680008	165.900955	165.904907	173.284042	-9.18	200.0	...
4	2021-10-01 16:00:00.723	1	200.0	173.820908	165.723297	165.900543	165.904861	173.280426	-9.18	200.0	...
5	2021-10-01 16:00:00.870	1	200.0	173.820908	165.806335	165.901611	165.904968	173.276825	-9.18	200.0	...
6	2021-10-01 16:00:01.023	1	200.0	173.824905	165.848724	165.902130	165.905014	173.273270	-9.18	200.0	...
7	2021-10-01 16:00:01.173	1	200.0	173.825897	165.911987	165.903549	165.905151	173.269730	-9.18	200.0	...
8	2021-10-01 16:00:01.323	1	200.0	173.821899	165.975845	165.905457	165.905350	173.266220	-9.18	200.0	...
9	2021-10-01 16:00:01.470	1	200.0	173.825897	166.084717	165.909119	165.905716	173.262741	-9.18	200.0	...

10 rows  $\times$  50 columns



# Bibliography

- [1] Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, “Time Series Management Systems: A Survey”, *TKDE*, 29(11), 2017.
- [2] Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, “ModelarDB: Modular Model-Based Time Series Management with Spark and Cassandra”, *PVLDB*, 11(11), 2018.
- [3] Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, “Demonstration of ModelarDB: Model-Based Management of Dimensional Time Series”, in *SIGMOD*, 2019.
- [4] Søren Kejser Jensen “Model-Based Time Series Management at Scale”, *PhD Thesis*, 2019.
- [5] Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, “Scalable Model-Based Management of Correlated Dimensional Time Series in ModelarDB<sub>+</sub>”, in *ICDE*, 2021.


Thank you!  
Abduvoris Abduvakhobov  
abduvorisa@cs.aau.dk



**AALBORG UNIVERSITY**  
DENMARK

# ModelarDB Architecture

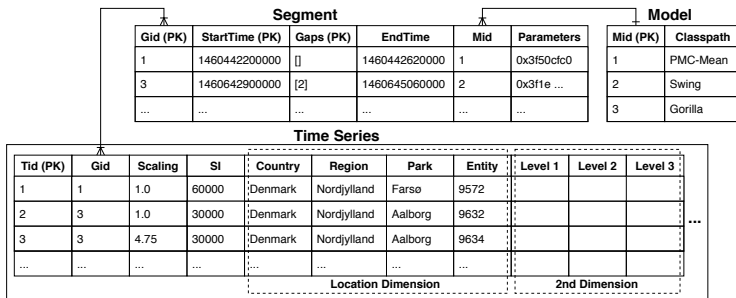
- ModelarDB is a portable Java library (ModelarDB Core) interfaced with a query engine (Apache Spark) and storage (Apache Cassandra).



- The architecture of a worker consists of three sets of components: Data Ingestion, Query Processing, and Segment Storage.



# ModelarDB Segment Structure



- Time Series and Model store metadata for time series and model types.
- Segment stores sub-sequences of time series as segments with a model.

# ModelarDB

## Model-Based Time Series Management at Scale

October 30<sup>th</sup>, 2020

Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen  
`skj, tbp, chr@cs.aau.dk`

Center for Data Intensive Systems, Daisy  
Department of Computer Science  
Aalborg University  
Denmark



**AALBORG UNIVERSITY**  
DENMARK



Center for Data-Intensive Systems

# Swing

## Implementation

- Models can always reconstruct their data points within the error bound:

```
@Override
protected float get(long timestamp, int index) {
    return (float) (this.a * timestamp + this.b);
}
```

- Many aggregates (e.g, SUM) can be computed directly from the models:

```
@Override
public double sum() {
    double first = this.a * this.getStartTime() + this.b;
    double last = this.a * this.getEndTime() + this.b;
    double average = (first + last) / 2;
    return average * this.length();
}
```



# Agenda

Motivation

Model-Based Compression

Implementation

Using ModelarDB

Performance

Summary and Open Issues



# Configuration

## Using ModelarDB

# Engine, Storage, and Query Interface

```
modelardb.engine spark://10.0.0.1:7077      # local,spark
modelardb.storage cassandra://10.0.0.1      # RDBMS, Cassandra, more
modelardb.interface http                    # none, socket, file, http
```

# Engine, Storage, and Query Interface

```
modelardb.source /home/user/DataSet/*.gz    # filepath, ip:port
modelardb.dimensions /home/user/dimensions.txt
modelardb.correlation auto                  # primitives
```

# Model Types

```
modelardb.model dk.aau.modelardb.core.models.PMC_MeanModel
modelardb.model dk.aau.modelardb.core.models.SwingFilterModel
modelardb.model dk.aau.modelardb.core.models.FacebookGorillaModel
```

# Error Bound, Length Limit Bound, and Latency Bound

```
modelardb.error 1      # per data point error bound in percentage
modelardb.limit 50     # length limit used by lossless model types
modelardb.latency 0    # maximum latency in terms of data points
```



# Deployment

## Using ModelarDB

- ModelarDB is simple to run standalone on a single node:

```
java -jar target/scala-2.11/ModelarDB-assembly-1.0.jar
```

- ModelarDB is simple to run on a Spark cluster by using spark-submit:

```
spark-submit
--conf spark.sql.parquet.filterPushdown=true
--conf spark.sql.orc.filterPushdown=true
--conf spark.cassandra.connection.host=10.0.0.2
--conf spark.hadoop dfs.replication=1
--driver-memory 4G
--executor-memory 3G
--master spark://10.0.0.1:7077
--jars UserDefinedModels
--class dk.aau.modelardb.Main
ModelarDB-assembly-1.0.jar
```

- Additional query engines can easily be interfaced with ModelarDB Core.



# Data Point View

## Using ModelarDB

```
# Client
curl -d 'SELECT MIN(val) FROM DataPoint WHERE sid = 17861' ...
{
  "time": "PT0.55S",
  "query": "SELECT MIN(val) FROM DataPoint WHERE sid = 17861",
  "result": [
    {"min(val)":48.0}
  ]
}

# Server
Data Point required columns{ val sid }
Data Point filters{ EqualTo(sid,17861) }
Segment required columns{ sid st et res mid param gaps }
Segment provided filters{ EqualTo(sid,17861) }
Segment rewritten filters{ EqualTo(gid,15408) }
Cache miss
Constructed predicates(gid=15408, takeWhile(st <= null, Gid IN )
Caching RDD
```



# Segment View

## Using ModelarDB

```
# Client
curl -d 'SELECT MIN_S(#) FROM Segment WHERE sid = 17861' ...
{
  "time": "PT0.295S",
  "query": "SELECT MIN_S(#) FROM Segment WHERE sid = 17861",
  "result": [
    {"maxs(sid, st, et, res, mid, param, gaps)":48.0}
  ]
}

# Server
Segment required columns { sid st et res mid param gaps }
Segment provided filters { EqualTo(sid,17861) }
Segment rewritten filters { EqualTo(gid,15408) }
Cache hit
```

- The Segment View has UDAFs for simple and aggregates over time.
- Multi-dimensional queries are supported with the use of GROUP BY.





# Agenda

Motivation

Model-Based Compression

Implementation

Using ModelarDB

Performance

Summary and Open Issues



# Environment

## Performance

- The evaluation uses real-life data sets from the energy domain.
- Most experiments are performed on a modest local cluster.
- Scalability experiments are performed using Microsoft Azure.
- ModelarDB configurations: with no grouping (MDB<sub>+</sub>-G), with auto (MDB<sub>+</sub>+GA), and with the best primitives per data set (MDB<sub>+</sub>+GB).
- For comparison the evaluation includes the state-of-the-art big data systems and file formats most commonly used in industry.<sup>1</sup>
  - InfluxDB, Apache Cassandra, Apache Parquet, and Apache ORC.
- No model-based time series management system is publicly available.
- Small, large, and multi-dimensional aggregates are evaluated using:
  - A Java Library (J), a Spark DataFrame (F), the Segment View (S), or the Data Point View (DP).

---

<sup>1</sup>From meeting with companies, our survey, and <https://db-engines.com/en/>



# Data Sets

## Performance

- EP** 45,353 time series collected from energy producers with a sampling interval of 60s and occupying 339 GiB as uncompressed CSV.
- It contains two dimensions with each containing two levels:
    - **Production:** *Entity*  $\rightarrow$  *Type*  $\rightarrow$   $\top$
    - **Measure:** *Concrete*  $\rightarrow$  *Category*  $\rightarrow$   $\top$
  - ActiveProductionMWh, AirPressure, DerateDetectionStage, EnergyInAccumulationTankMWh, EstimatedProductionMWh, EstimatedReductionMWh, EstimatedRunningCapacityMw, EstimatedSetPointPercentage, HorizontalIrradiance, Humidity, LinearScaledProductionMWh, NacelleDirection, PossibleProdMwh, Temperature, TiltedIrradiance, Tso, TsoOnshore, WindDirection, WindSpeedEastWest, WindSpeed, WindSpeedNorthSouth
  - $\text{MDB}_+ + \text{GA}$  time series with the same category are grouped per entity.
  - $\text{MDB}_+ + \text{GB}$  energy production measurements are grouped per entity.


# Data Sets

## Performance


- EF** 197 time series collected from wind turbines with a sampling interval of 200ms and occupying 372 GiB as uncompressed CSV.
- It contains two dimensions with three and two levels respectively:
    - **Location:**  $Entity \rightarrow Park \rightarrow Country \rightarrow \top$
    - **Measure:**  $Concrete \rightarrow Category \rightarrow \top$
  - AlmActCount, BecBulbSt, ConSt, DateTimeUTCDiff, ExTmp, GnTmpBrg1, GnTmpBrg2, GriPF, GriPhV1, GriPhV2, GriPhV3, Hz, PtAngValBl1, RotSpd, Spd, VAr, W, YwSt
  - MDB<sub>+</sub>+GA groups all time series on concrete measure and park.
  - MDB<sub>+</sub>+GB in general, group time series on the concrete measure and country, but time series are also grouped on the category and park.

# Ingestion and Storage Performance


- Results with groups of correlated time series are shown with stripes.



Ingestion Rate, EP



Storage Used, EP

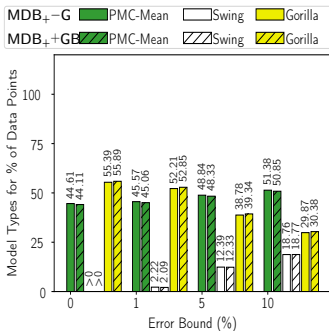


Storage Used, EH

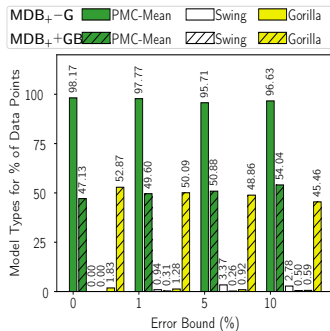
- Ingestion rate is 2.12–13.7 times faster than the industry formats.
- The storage required is reduced by 1.09–16.17 times for EP and by 3.12–112.64 times for EH compared to the industrial formats.
- Grouping series increases ingestion rate and can lower storage usage.

# Model Types

## Performance



Model Types, EP




Model Types, EH

- ModelarDB automatically adapts to each data set and error bound.
- Grouping increases ModelarDB's use of the lossless model type.




# Aggregate Queries


## Performance



Large Scale, EP



Small Scale, EP




Small Scale, EF

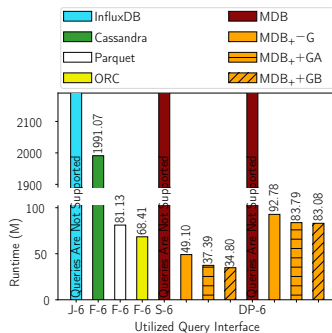
- Grouping time series decreases query time for large scale aggregate queries, but can increase query time for small scale aggregate queries.
- Experiments run on Azure show that ModelarDB scales linearly.

# Multi-Dimensional Queries EP

## Performance



Month and Category, EP




Month and Concrete, EP

- Grouping reduces query time as each query only reads whole groups.




# Multi-Dimensional Queries EF

## Performance



Month and Park, EF



Month and Entity, EF

- Grouping increases query time as the cluster is not fully utilized.